


x86-64 calling conventions:

Argument registers:	%rdi, %rsi, %rdx, %rcx, %r8, %r9
Preserved registers:	%rbp, %rbx, %r12, %r13, %r14, %r15
Return value:	%rax, %rdx



mov	

### Potpourri of déjà-vu multiple choices [6 points]

1. (2 points) How many times does function `execv()` return if there are no errors?  

<input type="radio"/> 0 times	<input type="radio"/> 2 times
<input type="radio"/> 1 time	<input type="radio"/> 3 times
  
2. (2 points) What is the difference between the `%rbx` and the `%ebx` register on an x86-64 machine?  

<input type="radio"/> <code>%ebx</code> refers to only the high order 32 bits of the <code>%rbx</code> register	<input type="radio"/> <code>%ebx</code> refers to only the low order 32 bits of the <code>%rbx</code> register
<input type="radio"/> they are totally different registers	<input type="radio"/> nothing, they are the same register
  
3. (2 points) What do you **NOT** expect to find in a Stack Frame?  

<input type="radio"/> Saved registers	<input type="radio"/> Function name
<input type="radio"/> Local (automatic) variables	<input type="radio"/> Return address

### Malloc [20 points]

1. (4 points) Given a system with a small amount of memory, explain which of the following memory allocation algorithm you would use: next-fit or best-fit. Make sure you explain the advantages and disadvantages of each of them.

2. (16 points) Consider an allocator implementation with the following characteristics:

- The **first-fit** free algorithm is used to allocate data.
- All blocks have a header with a size and a pointer to the previous block.
- The header is 16B (2\*8bytes) in size.
- Positive sizes indicate the block is allocated, and negative sizes indicate it is free.
- All freed blocks are immediately coalesced if possible.
- When a block is split, the lower (first) part of the block becomes the allocated part and the upper (second) part becomes the new free block.
- If the heap doesn't have enough space to hold the data, it grows by the minimum amount needed to fit the data. Always successfully.

For the given a heap representation, only the metadata is displayed. E.g., the following heap contains an allocated block of size 16, followed by a free block of size 32. The top row contains memory addresses, and the bottom row contains the values stored at those memory addresses.

Address	0xa000	0xa008	...	0xa020	0xa028	...
Value	16	0x0000	...	-32	0xa000	...

- a. (4 points) Assuming an **initially empty heap**, and given the current state of the heap represented below, which of the malloc sequence was executed?

Address	0xa000	0xa008	...
Value	-64	0x0000	...

- |  |  |
|--|--|
| <input type="radio"/> <code>p0 = malloc(32);</code><br><code>free(p0);</code><br><code>p0 = malloc(32);</code><br><code>free(p0);</code> | <input type="radio"/> <code>p0 = malloc(16);</code><br><code>p1 = malloc(32);</code><br><code>free(p0);</code><br><code>free(p1);</code> |
| <input type="radio"/> <code>p0 = malloc(64);</code>  | <input type="radio"/> <code>p0 = malloc(32);</code><br><code>free(p0);</code><br><code>p1 = malloc(16);</code><br><code>free(p1);</code> |

- b. (4 points) Assuming the heap starts as drawn in the previous question, and given the final state of the heap represented below, which of the malloc sequence was executed?

Address	0xa000	0xa008	...	0xa030	0xa038	...
Value	32	0x0000	...	-16	0xa000	...

- |  |  |
|--|--|
| <input type="radio"/> <code>p0 = malloc(32);</code><br><code>free(p0);</code><br><code>p0 = malloc(16);</code><br><code>free(p0);</code> | <input type="radio"/> <code>p0 = malloc(32);</code><br><code>free(p0);</code><br><code>p0 = malloc(16);</code> |
| <input type="radio"/> <code>p0 = malloc(32);</code>  | <input type="radio"/> <code>p0 = malloc(16);</code>  |

- c. (4 points) Assuming the heap starts as drawn above (b), if the following malloc executes, what is the value stored in p1?

`p1 = malloc(16);`

- |                              |                              |                              |                              |
|------------------------------|------------------------------|------------------------------|------------------------------|
| <input type="radio"/> 0xa010 | <input type="radio"/> 0xa018 | <input type="radio"/> 0xa020 | <input type="radio"/> 0xa028 |
| <input type="radio"/> 0xa030 | <input type="radio"/> 0xa038 | <input type="radio"/> 0xa040 | <input type="radio"/> 0xa048 |

- d. (4 points) Assuming the heap starts as drawn above (b), which value can fill the blank to successfully free the first block?

`free(_____);`

- |                              |                              |                              |                              |
|------------------------------|------------------------------|------------------------------|------------------------------|
| <input type="radio"/> 0xa000 | <input type="radio"/> 0xa008 | <input type="radio"/> 0xa010 | <input type="radio"/> 0xa018 |
| <input type="radio"/> 0xa020 | <input type="radio"/> 0xa028 | <input type="radio"/> 0xa030 | <input type="radio"/> 0xa038 |

**Assembly [25 points]****3. (9 points) Assembly and Reverse-Engineering**

Consider the following assembly dump

0000000000001139 &lt;bloop&gt;:

1139:	55	push	%rbp
113a:	48 89 e5	mov	%rsp,%rbp
113d:	48 83 ec 10	sub	\$0x10,%rsp
1141:	48 89 7d f8	mov	%rdi,-0x8(%rbp)
1145:	48 83 7d f8 29	cmpq	\$0x29,-0x8(%rbp)
114a:	7f 1b	jg	1167 <bloop+0x2e>
114c:	48 8b 05 dd 2e 00 00	mov	0x2edd(%rip),%rax
1153:	48 89 c6	mov	%rax,%rsi
1156:	48 8d 3d b5 0e 00 00	lea	0xeb5(%rip),%rdi
115d:	b8 00 00 00 00	mov	\$0x0,%eax
1162:	e8 c9 fe ff ff	callq	1030 <printf@plt>
1167:	90	nop	
1168:	c9	leaveq	
1169:	c3	retq	

**a. (2 points)** How many function arguments are defined in the above function bloop?
☐ 0
 ☐ 1
 ☐ 2
 ☐ 3
**b. (2 points)** How many local variables (not arguments) are declared in the above function bloop?
☐ 0
 ☐ 1
 ☐ 2
 ☐ 3
**c. (2 points)** Which of the following multiplies the value within the rax register by 9?
☐ lea (,rax,9), rax
 ☐ lea (rax,rax,8), rax  
☐ lea (rax,rax,9), rax
 ☐ lea 9(rax), rax
**d. (2 points)** If I saw `mov %rax, -0x8(%rbp)`, and given char is 1B, short 2B, int 4B, and long 8B, I would say that this local variable is what integer type?
☐ int
 ☐ long
 ☐ char
 ☐ short
**e. (1 points)** How many loops does the function above have?
☐ 0
 ☐ 1
 ☐ 2
 ☐ 3

4. (16 points) Reverse Engineering x86-64 to C. Consider the following x86-64 assembly and fill-in the blanks in the C code. Use the following variable-to-register mapping:

n <==> %rdi	y <==> %rsi	x <==> %ecx
result <==> %eax	i <==> %edx	

You may only use the symbolic variables such as `i`, `x`, and `result` in your C expressions  
 ---- do not use register names.

```

loop:
    movl    $10, %ecx
    movl    $0, %eax
    movl    $0, %edx
    jmp     .L2
.L3:
    leaq    (%rax,%rax,4), %rsi
    movq    %rcx, %rax
    leaq    (%rcx,%rsi), %rcx
    addq    $1, %rdx
.L2:
    leaq    (%rdi,%rdi), %rsi
    cmpq    %rdx, %rsi
    jg      .L3
    ret
  
```

```

long loop(long n)
{
    long i = _____;

    long result = _____;

    long x = _____;

    while (_____)
    {

        long y = _____;

        result = _____;

        x = _____;

        i++;
    }

    return _____;
}
  
```

## Buffer Overflow [15 points]

5. (15 points) Consider a program containing this poor-quality code, procedure vulnerable has the following disassembled form on a x86-64 machine:

<pre>void vulnerable(char t) {     char password[6];     char name[4];     gets(name);     password[0]='H';     password[1]='e';     password[2]='l';     password[3]='l';     password[4]='o';     password[5]=t;     printf("You cannot know my password %s!\n", name);     // here }</pre>	<pre>vulnerable: # @vulnerable     pushq %rbp     movq %rsp, %rbp     subq \$0x10, %rsp     movb %dil, -1(%rbp) # dil→8lsb of rdi     leaq -0xb(%rbp), %rdi     callq gets     movb \$72, -7(%rbp) # H     movb \$101, -6(%rbp) # e     movb \$108, -5(%rbp) # l     movb \$108, -4(%rbp) # l     movb \$111, -3(%rbp) # o     movb -1(%rbp), %al     movb %al, -2(%rbp)     leaq -0xb(%rbp), %rsi     mov \$0x400104, %rdi # "You cannot ..."     callq printf     addq \$0x10, %rsp     popq %rbp     retq</pre>
---	--

For the following questions, recall that:

- gets is the standard C library routine.
- x86-64 machines are **little-endian**.
- C strings are null-terminated (i.e., terminated by a character with value 0x00).

Consider the case where procedure vulnerable is called with argument t equal to '\0', and we type "Luis" in response to gets.

- a. (3 points) Which elements of array password were overwritten when gets is called?

- |   |   |
|---|---|
| <input type="radio"/> None                | <input type="radio"/> Only password[0]          |
| <input type="radio"/> password[0] and [1] | <input type="radio"/> password[0], [1], and [2] |

- b. (3 points) Which of the following stack values were corrupted?

- |   |  |
|---|--|
| <input type="radio"/> Arguments                     | <input type="radio"/> Saved registers            |
| <input type="radio"/> Arguments and Saved registers | <input type="radio"/> None of the listed options |

- c. **(3 points)** What (exactly) was printed by the function when printf was executed?

Answer:

- d. **(3 points)** What would **NOT** help with preventing code injection attacks?

- |   |   |
|---|---|
| <input type="radio"/> Add a canary to the stack     | <input type="radio"/> Make the stack larger than it needs to be |
| <input type="radio"/> Make the stack not executable | <input type="radio"/> Randomize the memory address of the stack |

- e. **(3 points)** Which of the following cases represents a buffer overflow?

- |   |   |
|---|---|
| <input type="radio"/> Writing a string of length 5 to an array of chars of size 6 | <input type="radio"/> Reading 20B from an array of size 10B |
| <input type="radio"/> Writing 5B of code into a stack buffer with 10B of capacity | <input type="radio"/> Executing code in the stack           |



**Creation and execution of programs [24 points]**

---

6. **(6 points)** Compiling, linking, loading. For each question, determine at which stage the action happens.

a. (2 points) Determine the location of functions in other object files?

☐ Compiling

☐ Loading

☐ Linking

☐ Assembling

b. (2 points) Place shared libraries in memory?

☐ Compiling

☐ Loading

☐ Linking

☐ Assembling

c. (2 points) Determine the offset of a stack variable?

☐ Compiling

☐ Loading

☐ Linking

☐ Assembling

7. **(9 points)** Answer the questions for the code below.

```
static int x = 0, y = 5;

int what_is_this(void) {
    x = x + y;
    y = y + 1;
    return y;
}

int main(void) {
    int v = what_is_this();
    printf("%d\n", v);
    return v;
}
```

- a. **(3 points)** With respect to the Linker, which of the following is a **global** symbol?

☐ what\_is\_this

☐ y

☐ x

☐ v

- b. **(3 points)** With respect to the Linker, which of the following is a **local** symbol?

☐ v

☐ what\_is\_this

☐ x

☐ main

- c. (3 points) With respect to the Linker, which of the following is **NOT** registered as a symbol?

☐ v                      ☐ what\_is\_this                      ☐ x                      ☐ main

8. (9 points) For each question, read the code – assume all code is given!

- a. (3 points) What type of error will the code produce?

```
int main() {  
    int a[4] = {0};  
    for (int i=0 ; i<10000 ; i++) {  
        a[i] = i;  
    }  
    return 0;  
}
```

☐ Runtime error                      ☐ Compilation error  
☐ Linking error                      ☐ Loading error

- b. (3 points) What type of error will the code produce?

```
int main() {  
    int a[4] = {0};  
    for (int i=0 ; i<4 ; i++) {  
        a[i] = add(a[1], i);  
    }  
    return 0;  
}
```

☐ Runtime error                      ☐ Compilation error  
☐ Linking error                      ☐ Loading error

- c. (3 points) What type of error will the code produce?

```
int add(int a, int b);  
int main() {  
    int a[4] = {0};  
    for (int i=0 ; i<4 ; i++) {  
        a[i] = add(a[1], i);  
    }  
    return 0;  
}
```

☐ Runtime error                      ☐ Compilation error  
☐ Linking error                      ☐ Loading error